

```

In [12]: import pandas as pd
from sklearn.compose import ColumnTransformer
from sklearn.neural_network import MLPRegressor
from sklearn.metrics import mean_squared_error
import seaborn as sns
import matplotlib.pyplot as plt

df = pd.read_csv(r"C:\Users\adamg\Documents\ML_Research_Final_CSV.csv", header=0)
print(f"Read in {len(df)} rows")
df1 = pd.read_csv(r"C:\Users\adamg\Documents\ML_Research_Final_CSV.csv")#, header=0
df.head()
#C:\Users\adamg\Documents\ML_Research_Final_CSV.csv
df.replace("?", 10000, inplace=True) #10,000 is way beyond the range of column values
#print(df)
fig1=sns.pairplot(data=df1)
plt.show()
df.drop([0], 0, inplace=True)
print('Data:',df)
df.head()

import numpy as np
from sklearn.model_selection import train_test_split

X_1 = np.array(df.drop([0], 1)) #Last column contains label, so ignore it when creating features
#print(f"{X}")
y_1 = np.array(df[0]) #second column is a label which is our y(Pressure)
# Using the Kaggle Scaling

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import pickle
import timeit
start = timeit.default_timer()
clf = MLPRegressor(solver='lbfgs', alpha=1e-5,
                  hidden_layer_sizes=(5, 2), random_state=51)
#solver adam works for this as well
X_train, X_test, y_train, y_test = train_test_split(X_1, y_1, test_size=0.4, random_state=51)
clf.fit(X_train,y_train)
clf.score(X_test,y_test)
print(f"Accuracy of MLP Regressor is:{clf.score}")
print("clf = ",clf.score)
print("X_train = ",X_train.shape)
print("X_test = ",X_test.shape)
print("y_train = ",y_train.shape)
print("y_test = ",y_test.shape)
# WORKS!
#X_train, X_test, y_train, y_test = train_test_split(X_1, y_1, test_size=0.25, random_state=51)
nn_model = MLPRegressor(solver='adam', alpha=1e-5, hidden_layer_sizes=(100,100), random_state=51)
#lbfgs
#y_train=y_train.astype('int')
#y_test=y_test.astype('int')
#X_test=X_test.astype('int')
#X_train=X_train.astype('int')
nn_model.fit(X_train, y_train)
nn_accuracy = nn_model.score(X_test, y_test)# Why is accuracy not used here

```

```

stop = timeit.default_timer()
prediction = nn_model.predict(X_test)
prediction=prediction.astype('float')
y_test=y_test.astype('float')
mse = mean_squared_error(y_test, prediction)
#print('Predict acc:', prediction-y_test)
print('Time:', stop-start)
print(f"Mean Squared Error is :{mse}")
#Use the mean of the y_test and prediction for mse because
#if the numbers are simply small, you can still get a low mse
# import mean squared error


nn_model.fit(X_train, y_train) #doubt
plt.plot(nn_model.loss_curve_, label="train") #doubt
plt.plot(nn_model.loss_curve_, label="test")
plt.legend()
plt.xlabel('epoch')
plt.ylabel('Loss')

```

Time: 0.1251195999997251

Mean Squared Error is :0.20253208920140628

Out[12]: Text(0, 0.5, 'Loss')



Mse: v_x = 0.05228664057237373

Mse: v_y = 0.018790183701411427

In []: